

Installation Guide - Crunchy Containers for PostgreSQL

REVISION HISTORY			
NUMBER	DATE	DESCRIPTION	NAME

Contents

1	Project Setup & Docker Installation	1
1.1	Assumptions	1
1.2	Step 1 - Project Directory Structure	1
1.3	Step 2 - Install the Host Dependencies	2
1.4	Step 3 - Build the Containers	2
2	OpenShift Environment	3
2.1	Installation	3
2.2	Configure NFS for Persistence	3
2.3	Examples	4
3	Kubernetes Environment	5
3.1	Installation	5
3.2	Examples	5
3.3	Tips	5
4	Legal Notices	6

1 Project Setup & Docker Installation

The crunchy-containers can run on different environments including:

- Standalone Docker
- OpenShift Enterprise
- Kubernetes 1.2.X

In this document we list the basic installation steps required for these environments.

1.1 Assumptions

The install assumes the following about your host configuration:

- Centos 7 or RHEL 7 VM
- user called **someuser** has been created
- **someuser** has sudo privileges with no password

1.2 Step 1 - Project Directory Structure

First add the following lines to your .bashrc file to set the project paths:

```
export GOPATH=$HOME/cdev
export GOBIN=$GOPATH/bin
export PATH=$PATH:$GOBIN
export CCP_BASEOS=centos7
export CCP_PGVERSION=9.6
export CCP_VERSION=1.5
export CCP_IMAGE_TAG=$CCP_BASEOS-$CCP_PGVERSION-$CCP_VERSION
export CCROOT=$GOPATH/src/github.com/crunchydata/crunchy-containers
```

You will then need to log out and back in for the changes to your .bashrc file to take effect.

Next, set up a project directory structure and pull down the project:

```
mkdir $HOME/cdev $HOME/cdev/src $HOME/cdev/pkg $HOME/cdev/bin
cd $GOPATH
sudo yum -y install golang git docker postgresql
go get github.com/tools/godep
cd src/github.com
mkdir crunchydata
cd crunchydata
git clone https://github.com/crunchydata/crunchy-containers
cd crunchy-containers
git checkout 1.5
godep restore
```

If you are a Crunchy enterprise customer, you will place the CRUNCHY repo key and yum repo file into the \$CCROOT/-conf directory at this point. These files can be obtained through <https://access.crunchydata.com/> on the downloads page.

1.3 Step 2 - Install the Host Dependencies

Next, install system dependencies:

```
sudo yum -y update
sudo groupadd docker
sudo usermod -a -G docker someuser
```

Remember to log out of the **someuser** account for the Docker group to be added to your current session. Once it's added, you'll be able to run Docker commands from your user account.

```
su - someuser
```

You can ensure your **someuser** account is added correctly by running the following command and ensuring **docker** appears as one of the results:

```
groups
```

Before you start Docker, you might consider configuring Docker storage: This is described if you run:

```
man docker-storage-setup
```

Follow the instructions available [on the main OpenShift documentation page](#) to configure Docker storage appropriately.

These steps are illustrative of a typical process for setting up Docker storage. You will need to run these commands as root.

First, add an extra virtual hard disk to your virtual machine (see [this blog post](#) for tips on how to do so).

Run this command to format the drive, where `/dev/sd?` is the new hard drive that was added:

```
fdisk /dev/sd?
```

Next, create a volume group on the new drive partition within the `fdisk` utility:

```
vgcreate docker-vg /dev/sd?
```

Then, you'll need to edit the `docker-storage-setup` configuration file in order to override default options. Add these two lines to **/etc/sysconfig/docker-storage-setup**:

```
DEVS=/dev/sd?
VG=docker-vg
```

Finally, run the command **docker-storage-setup** to use that new volume group. The results should state that the physical volume `/dev/sd?` and the volume group `docker-vg` have both been successfully created.

Next, we enable and start up Docker:

```
sudo systemctl enable docker.service
sudo systemctl start docker.service
```

1.4 Step 3 - Build the Containers

At this point, you have a decision to make - either download prebuilt containers from [Dockerhub](#), **or** build the containers on your local host.

To download the prebuilt containers, make sure you can login to [Dockerhub](#), and then run the following:

```
docker login
cd $CCPROOT
./bin/pull-from-dockerhub.sh
```

Or if you'd rather build the containers from source, perform a container build as follows:

```
cd $CCPROOT
make setup
make all
```

After this, you will have all the Crunchy containers built and are ready for use in a **standalone Docker** environment.

2 OpenShift Environment

2.1 Installation

See the OSE installation guide for details on how to install OSE on your host. The main instructions are here:

https://docs.openshift.com/enterprise/3.2/install_config/install/index.html

Or, if you'd prefer to install OpenShift Origin, the easiest way to get OpenShift Origin up and running is found here: https://github.com/openshift/origin/blob/master/docs/cluster_up_down.md

For examples and tips on how to run OpenShift Enterprise & Origin, please look at the `examples.adoc` documentation.

2.2 Configure NFS for Persistence

NFS is required for some of the OpenShift examples, those dealing with backups and restores will require a working NFS for example.

First, if you are running your NFS system with SELinux in enforcing mode, you will need to run the following command to allow NFS write permissions:

```
sudo setsebool -P virt_use_nfs 1
```

The necessary dependencies will need to be installed and started next:

```
sudo setsebool -P virt_use_nfs 1
sudo yum -y install nfs-utils libnfsidmap
sudo systemctl enable rpcbind nfs-server
sudo systemctl start rpcbind nfs-server rpc-statd nfs-idmapd
```

Next, you will need to set the permissions of your NFS path so that your pods can have write access. For the Crunchy examples, the **nfsnobody** GUI was chosen as an example. Pods will reference the **nfsnobody** GID (65534) as a security context **supplementalGroup** attribute. This setting will allow the pod to have group permissions of 65534 and therefore be able to write to the NFS persistent volumes.

The permissions on the NFS path are set as follows:

```
sudo mkdir /nfsfileshare
sudo chmod 777 /nfsfileshare
```

Most of the Crunchy containers run as the postgres UID (26), but you will notice that when **supplementalGroups** are specified, the pod will include the **nfsnobody** group in the list of groups for the pod user.

The case of Amazon file systems is different, for that you use the **fsGroup** security context setting but the idea for allowing write permissions is the same.

You will then need to modify the `/etc/exports` file.

```
sudo vi /etc/exports
```

The `/etc/exports` file should contain a line similar to this one except with the applicable IP address specified:

```
/nfsfileshare 192.168.122.9(rw, sync)
```

After saving and closing the file, you will need to re-export all directories:

```
sudo exportfs -r
```

Detailed instructions that you can use for setting up a NFS server on Centos 7 are provided in the following link.

<http://www.itzgeek.com/how-tos/linux/centos-how-tos/how-to-setup-nfs-server-on-centos-7-rhel-7-fedora-22.html>

OpenShift NFS examples can be found here:

<https://github.com/openshift/origin/tree/master/examples/wordpress/nfs>

The examples specify a test NFS server running at IP address 192.168.0.103.

On that server, the /etc/exports file looks like this:

```
/nfsfileshare *(rw, sync)
```

Test your NFS configuration out by mounting a local directory:

```
mount 192.168.0.114:/nfsfileshare /mnt/nfsfileshare
```

if you are running your client on a VM, you will need to add *insecure* to the exportfs file on the NFS server, this is because of the way port translation is done between the VM host and the VM instance.

For more details on this bug, please see the following link.

<http://serverfault.com/questions/107546/mount-nfs-access-denied-by-server-while-mounting>

A suggested best practice for tuning NFS for PostgreSQL is to configure the PostgreSQL fstab mount options like so:

```
proto=tcp, suid, rw, vers=3, proto=tcp, timeo=600, retrans=2, hard, fg, rsize=8192, wsize ↵  
=8192
```

Network options:

```
MTU=9000
```

If interested in mounting the same NFS share multiple times on the same mount point, look into the [noac mount option](#).

2.3 Examples

For running the examples that require persistent volumes, you will need to run the following script:

```
cd $CCPROOT/examples/pv  
./create-pv.sh  
./create-pvc.sh
```

View the README.txt for command-line usage.

If you are wanting to run the examples on a Minishift instance you will need to create the PVs using hostPath as follows:

```
oc login -u system:admin  
./create-pv.sh hostpath  
oc login -u developer  
./create-pvc.sh
```

Additional steps are required to allow persistence to work on Minishift including:

```
oc login -u system:admin
oc edit scc restricted
```

Above, you will change `runAsUser.Type` strategy to `RunAsAny`.

On the `boot2docker` instance running `Minishift`, you will need to set the host path permissions as follows:

```
chmod 777 /mnt/sda1/data
```

The `NAMESPACE` environment variable is set to indicate which OpenShift project you want various example objects to use. This variable is set to **default** within the **examples/envvars.sh** script. Set this to match your project configuration.

See [here](#) to view the documentation showing various examples.

3 Kubernetes Environment

3.1 Installation

I recommend using `kubeadm` or `minikube` to try the examples out.

See the following links for installation instructions:

- <https://github.com/Kubernetes/minikube>
- <http://linode.com/containers/setup-kubernetes-kubeadm-centos/>
- <https://kubernetes.io/docs/getting-started-guides/kubeadm/>

3.2 Examples

The namespace is set for the examples within the **examples/envvars.sh** script and defaults to **default**. Set this variable according to your project configuration.

Note, some of the examples assume an NFS file system for creating persistent volumes. See above for details on setting NFS permissions and the use of **supplementalGroups** within pod specs.

Visit the [examples documentation](#) for different use cases and examples.

3.3 Tips

Make sure your hostname resolves to a single IP address in your `/etc/hosts` file! If not, the NFS examples will not work.

You should see a single IP address returned from this command:

```
hostname --ip-address
```

```
sudo PATH=$PATH ALLOW_PRIVILEGED=true ./hack/local-up-cluster.sh
```

Note: specifying `ALLOW_PRIVILEGED=true` is required if you are running in SELinux enforcing mode. This allows you to specify the following in your pod spec to run the container as privileged:

```
"securityContext": {
  "privileged": true
},
```


4 Legal Notices

Copyright © 2017 Crunchy Data Solutions, Inc.

CRUNCHY DATA SOLUTIONS, INC. PROVIDES THIS GUIDE "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF NON INFRINGEMENT, MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Crunchy, Crunchy Data Solutions, Inc. and the Crunchy Hippo Logo are trademarks of Crunchy Data Solutions, Inc.