



Docker Examples - Crunchy Containers for PostgreSQL

Crunchy Data Solutions, Inc.

Version 1.4, 2017-04-25

Docker Examples

Examples of running the Crunchy containers in a pure Docker environment are found in the `examples/docker` directory. Those examples are explained below. The examples make use of the local Docker volume driver and data volumes created by the **`docker create volume`** command. There is a bug however that requires the container to run as privileged when in selinux enforcing mode. See [here](#) for more details. Also, see [here](#). When this is fixed, you will not need to specify **privileged** when starting the examples.

basic - Running a single database container

Create the container with this command:

```
cd examples/docker/basic
./run.sh
```

This script will do the following:

- create a data volume to persist PostgreSQL data files to
- start up a container named **basic**
- initialize the database using the passed in environment variables, building a user, database, schema, and table based on the environment variable values.
- maps the PostgreSQL port of 5432 in the container to your local host port of 12000.

The container creates a default database called **testdb**, a default user called **testuser** with a default password of **testpsw**, you can use this to connect from your local host as follows:

```
psql -h localhost -p 12000 -U testuser -W testdb
```

To shut down the instance, run the following commands:

```
docker stop basic
```

To start the instance, run the following commands:

```
docker start basic
```

master-replica - Creating a master and replica database cluster

Run the container with this command:

```
cd examples/docker/master-replica
./run.sh
```

This script will do the following:

- create a docker volume using the local driver for the master
- create a docker volume using the local driver for the replica
- start up a container named master binding to port 12000
- start up a container named replica binding to port 12002
- initialize the database using the passed in environment variables, building a user, database, schema, and table based on the environment variable values.
- maps the PostgreSQL port of 5432 in the container to your local host port of 12000.

The container creates a default database called **testdb**, a default user called **testuser** with a default password of **testpsw**, you can use this to connect from your local host as follows:

```
psql -h localhost -p 12007 -U testuser -W userdb
psql -h localhost -p 12008 -U testuser -W userdb
```

To shut down the instance, run the following commands:

```
docker stop master replica
```

To start the instance, run the following commands:

```
docker start master replica
```

backup - Performing a backup

In order to run this backup script, you first need to edit run-backup.sh to specify your host IP address you are running on. The script assumes you are going to backup the container created in Example 2.

Run the backup with this command:

```
cd examples/docker/backup
./run-backup.sh
```

This script will do the following:

- start up a backup container named masterbackup

- run `pg_basebackup` on the container named master
- store the backup in `/tmp/backups/master` directory
- exit after the backup

restore - Performing a restore

In order to run this backup script, you first need to edit `run-restore.sh` to specify your host IP address you are running on. The script assumes you are going to backup the container created in Example 2.

Run the backup with this command:

```
cd examples/docker/restore
./run.sh
```

This script will do the following:

- start up a container named `masterrestore`
- copy the backup files from Example 3 into `/pgdata`
- start up the container using the backup files
- maps the PostgreSQL port of 5432 in the container to your local host port of 12001 as to not conflict with the master running in the previous example.

master-replica - Running PostgreSQL in Master-Slave Configuration

The container can be passed environment variables that will cause it to assume a PostgreSQL replication configuration with a master instance streaming to a slave replica instance.

The following env variables are specified for this configuration option:

- `PG_MASTER_USER` - The username used for master-slave replication value=`master`
- `PG_MASTER_PASSWORD` - The password for the PG master user
- `PG_USER` - The username that clients will use to connect to PG server value=`user`
- `PG_PASSWORD` - The password for the PG master user
- `PG_DATABASE` - The name of the database that will be created value=`userdb`
- `PG_ROOT_PASSWORD` - The password for the PG admin

This examples assumes you have run Example 1, and that the master container is running.

For running the master-slave configuration , you can run the following scripts:

```
run.sh
```

You can verify that replication is working by connecting to the replica as follows:

```
psql -h 127.0.0.1 -p 12002 -U postgres postgres
```

If you have created tables or data in the master database, they should show up in this replicated copy of that database.

crunchy-proxy - crunchy-proxy

This example shows how to use the **crunchy-proxy** to act as a smart proxy to a PostgreSQL cluster. The example depends upon the **master-replica** example being run prior.

crunchy-proxy offers a high performance alternative to pgbouncer and pgpool.

The proxy example copies a configuration file to the PV_PATH and starts up the **crunchy-proxy** within a Deployment.

The proxy reads the configuration file from a **/config** volume mount and begins execution.

Start by running the proxy container:

```
cd $CCPROOT/examples/docker/crunchy-proxy  
./run.sh
```

The proxy will listen on port 5432 as specified in the configuration file. The example creates a Service named **crunchy-proxy** that you can use to access the configured PostgreSQL backend containers from the **master-replica** example.

See the following link for more information on the **crunchy-proxy**:

<https://github.com/CrunchyData/crunchy-proxy>

Test the proxy by running psql commands via the proxy connection:

```
psql -h crunchy-proxy -U postgres postgres
```

SQL "reads" will be sent to the PostgreSQL replica database if your SQL includes the **crunchy-proxy** read annotation. SQL statements that do not include the read annotation will be sent to the master database container within the PostgreSQL cluster.

crunchy-proxy - a smart proxy for Postgres

A **crunchy-proxy** example is provided that will run a container that is configured to be used with the master and replica example provided in the **master-replica** example.

You can create the proxy by running:

```
cd $CCPROOT/examples/docker/crunchy-proxy
./run.sh
```

This proxy will listen on localhost:12432. You can access the **master-replica** cluster by:

```
psql -h localhost -p 12432 -U postgres postgres
```

See this link for details on the **crunchy-proxy**: <https://github.com/CrunchyData/crunchy-proxy>

You might consider **crunchy-proxy** over pgpool and pgbouncer if you need load-balancing and smart SQL routing.

pgpool - pgpool

A pgpool example is provided that will run a pgpool container that is configured to be used with the master and replica example provided in the **master-replica** example. After running those commands to create a master and replica, you can create a pgpool container by running the following example command:

```
cd examples/docker/pgpool
./run.sh
```

Enter the following command to connect to the pgpool that is mapped to your local port 12003:

```
psql -h localhost -U testuser -p 12003 userdb
```

You will enter the password of **password** when prompted. At this point you can execute both INSERT and SELECT statements on the pgpool connection. Pgpool will direct INSERT statements to the master and SELECT statements will be sent round-robin to both master and replica.

badger - pgbadger

A pgbadger example is provided that will run a HTTP server that when invoked, will generate a pgbadger report on a given database.

pgbadger reads the log files from a database to product an HTML report that shows various Postgres statistics and graphs.

To run the example, modify the run-badger.sh script to refer to the Docker container that you want to run pgbadger against, also referring to the container's data directory, then run the example as follows:

```
cd examples/docker/badger
./run.sh
```

After execution, the container will run and provide a simple HTTP command you can browse to view the report. As you run queries against the database, you can invoke this URL to generate updated reports:

```
curl http://127.0.0.1:14000/api/badgergenerate
```

dns - dns

Some users will need or want a DNS name to resolve to their container names. The crunchy-dns container provides the following:

- * listens to a Docker URL or socket for events that would indicate a container is created or destroyed
- * implements the consul.io DNS server
- * registers new container information into the DNS server
- * deregisters container information from the DNS server when a container is destroyed

Start the crunchy-dns server by running its container as follows:

```
cd examples/docker/dns
./run.sh
```

This is a privileged container and will bind to your local IP address at port 53. At this point, you can now start a new Postgres container and you should be able to do a DNS lookup as follows:

```
dig @192.168.122.138 containername.service.dc1.crunchy.lab
```

In this example, the local IP address of the DNS container is 192.168.122.138 and it assumes you have started a container named containername

You can alter the DNS domain name within the startup script if desired.

You can also browse the consul web UI at:

```
http://<your ip address>:8500/ui
```

metrics - metrics collection

You can collect various Postgres metrics from your database container by running a crunchy-collect

container that points to your database container.

Metrics collection requires you run the crunchy 'scope' set of containers that includes:

- prometheus
- prometheus push gateway
- grafana

To start this set of containers, run the following:

```
cd examples/docker/metrics
./run.sh
```

These metrics are described in this [document](#).

An example has been provided that runs a database container and also the associated metrics collection container, run the example as follows:

```
cd examples/docker/collect
./run.sh
```

Every 3 minutes the collection container will collect postgres metrics and push them to the crunchy prometheus database. You can graph them using the crunchy grafana container.

vacuum - vacuum

You can perform a postgres vacuum command by running the crunchy-vacuum container. You specify a database to vacuum using environment variables.

An example is shown in the examples/docker/run-vacuum.sh script and can be run as follows:

```
cd examples/docker/master-replica
./run.sh
```

This example performs a vacuum on a single table in the master postgres database. Vacuum is controlled via the following environment variables:

- VAC_FULL - when set to true adds the FULL parameter to the VACUUM command
- VAC_TABLE - when set, allows you to specify a single table to vacuum, when not specified, the entire database tables are vacuumed
- JOB_HOST - required variable is the postgres host we connect to
- PG_USER - required variable is the postgres user we connect with
- PG_DATABASE - required variable is the postgres database we connect to
- PG_PASSWORD - required variable is the postgres user password we connect with

- PG_PORT - allows you to override the default value of 5432
- VAC_ANALYZE - when set to true adds the ANALYZE parameter to the VACUUM command
- VAC_VERBOSE - when set to true adds the VERBOSE parameter to the VACUUM command
- VAC_FREEZE - when set to true adds the FREEZE parameter to the VACUUM command

custom-setup- custom setup.sql

You can use your own version of the setup.sql SQL file to customize the initialization of database data and objects when the container and database are created.

An example is shown in the examples/docker/custom-setup/run.sh script and can be run as follows:

```
cd examples/docker/custom-setup
./run.sh
```

This works by placing a file named, setup.sql, within the /pgconf mounted volume directory. Portions of the setup.sql file are required for the crunchy container to work, see comments within the sample setup.sql file.

pgbouncer - pgbouncer

The pgbouncer utility can be used to provide a connection pool to postgres databases. The crunchy-pgbouncer container also contains logic that lets it perform a failover from a master to a slave database.

To test this failover, you first create a running master/slave cluster as follows:

```
cd examples/docker/master-replica
./run.sh
```

An example is shown in the examples/docker/pgbouncer/run.sh script and can be run as follows:

```
cd examples/docker/pgbouncer
./run.sh
```

This example configures pgbouncer to provide connection pooling for the master and pg-replica databases. It also sets the FAILOVER environment variable which will cause a failover to be triggered if the master database can not be reached.

To trigger the failover, stop the master database:

```
docker stop master
```

At this point, the pgbouncer will notice that the master is not reachable and touch the trigger file on the configured slave database to start the failover. The pgbouncer container will then reconfigure pgbouncer to relabel the slave database into the master database so clients to pgbouncer will be able to connect to the master as before the failover.

To just log into the database from the pgbouncer connection pool you would enter the following:

```
psql -h localhost -p 12005 -U testuser master
```

sync - synchronous replication

This example, `examples/docker/sync`, provides a streaming replication configuration that includes both synchronous and asynchronous slaves.

To run this example, run the following:

```
cd examples/docker/sync
./run.sh
```

You can test the replication status on the master by using the following command:

```
psql -h 127.0.0.1 -p 12000 -U postgres postgres -c 'table pg_stat_replication'
```

You should see 2 rows, 1 for the async slave and 1 for the sync slave. The `sync_state` column shows values of `async` or `sync`.

You can test replication to the slaves by entering some data on the master like this, and then querying the slaves for that data:

```
psql -h 127.0.0.1 -p 12000 -U postgres postgres -c 'create table foo (id int)'
psql -h 127.0.0.1 -p 12000 -U postgres postgres -c 'insert into foo values (1)'
psql -h 127.0.0.1 -p 12002 -U postgres postgres -c 'table foo'
psql -h 127.0.0.1 -p 12003 -U postgres postgres -c 'table foo'
```

pgadmin - pgadmin4

This example, `examples/docker/pgadmin4`, provides a container that runs the pgadmin4 web application.

To run this example, run the following:

```
cd examples/docker/pgadmin4
./run.sh
```

You should now be able to browse to <http://YOURLOCALIP:5050> and log into the web application using a user ID of **admin@admin.org** and password of **password**. Replace YOURLOCALIP with whatever your local IP address happens to be.

pitrr - PITR (point in time recovery)

This example, `examples/docker/pitr`, provides an example of performing a point in time recovery.

To run this example, run the following to create a database container:

```
cd ./examples/docker/pitr
./run-master-pitr.sh
```

It takes about 1 minute for the database to become ready for use after initially starting.

This database is created with the `ARCHIVE_MODE` and `ARCHIVE_TIMEOUT` environment variables set. See the `pitr.asciidoc` for more details on these settings. Warning: this example writes the WAL segment files to the `/tmp` directory...running it for a long time could fill up your `/tmp`!

Next, we will create a base backup of that database using this:

```
./run-master-pitr-backup.sh
```

At this point, WAL segment files are created every 60 seconds that contain any database changes. These segments are stored in the `/tmp/master-data/master-wal` directory.

Next, create some data in your database using this command:

```
psql -h 127.0.0.1 -p 12000 -U postgres postgres -c "select
pg_create_restore_point('beforechanges')"
psql -h 127.0.0.1 -p 12000 -U postgres postgres -c 'create table pitrtest (id int)'
psql -h 127.0.0.1 -p 12000 -U postgres postgres -c "select
pg_create_restore_point('afterchanges')"
psql -h 127.0.0.1 -p 12000 -U postgres postgres -c "select
pg_create_restore_point('nomorechanges')"
psql -h 127.0.0.1 -p 12000 -U postgres postgres -c "checkpoint"
```

Next, stop the database to avoid conflicts with the WAL files while attempting to do a restore from them:

```
docker stop master-pitr
```

The commands above set restore point labels which we can use to mark the points in the recovery process we want to reference when creating our restored database. Points before and after the test table were made.

Next, let's edit the restore script to use the base backup files created in the step above. You can view the backup path name under `/tmp/backups/master-pitr` directory. You will see a value like **2016-09-21-21-03-29**. Copy and paste that value into the `run-restore-pitr.sh` script in the **BACKUP** environment variable.

```
vi ./run-restore-pitr.sh
```

Next, let's see if we can restore the database before we created the test table in the last command, we will restore using the backup and will use the **beforechanges** label as the restore target name in the PITR:

```
./run-restore-pitr.sh
```

The WAL segments are read and applied when restoring from the database backup. At this point, you should be able to verify that the database was restored to the point before creating the test table:

```
psql -h 127.0.0.1 -p 12001 -U postgres postgres -c 'table pitrtest'
```

This sql command should show that the `pitrtest` table does not exist at this recovery time.

PostgreSQL allows you to pause the recovery process if the target name or time is specified. This pause would allow a DBA a chance to review the recovery time/name and see if this is what they want or expect. If so, the DBA can run the following command to resume and complete the recovery:

```
psql -h 127.0.0.1 -p 12001 -U postgres postgres -c 'select pg_xlog_replay_resume()'
```

Until you run the statement above, the database will be left in read-only mode.

Next, run the script to restore the database to the **afterchanges** restore point, do this by updating the `RECOVERY_TARGET_NAME` to **afterchanges**:

```
vi ./run-restore-pitr.sh  
./run-restore-pitr.sh
```

After this restore, you should be able to see the test table:

```
psql -h 127.0.0.1 -p 12001 -U postgres postgres -c 'table pitrtest'  
psql -h 127.0.0.1 -p 12001 -U postgres postgres -c 'select pg_xlog_replay_resume()'
```

Lastly, let's recovery using all of the WAL files, this will get the restored database as current as possible, edit the script to remove the `RECOVERY_TARGET_NAME` environment setting completely:

```
./run-restore-pitr.sh
sleep 30
psql -h 127.0.0.1 -p 12001 -U postgres postgres -c 'table pitrtest'
psql -h 127.0.0.1 -p 12001 -U postgres postgres -c 'create table foo (id int)'
```

At this point, you should be able to create new data in the restored database and the test table should be present. When you recover the entire WAL history, resuming the recovery is not necessary to enable writes.

Other options exist for performing a PITR, see the `pitr.asciidoc` for full details.

pgaudit - pgaudit

This example, `examples/docker/pgaudit`, provides an example of enabling pgaudit output. As of release 1.3, pgaudit is included in the `crunchy-postgres` container and is added to the postgres shared library list in the `postgresql.conf`.

Given the numerous ways pgaudit can be configured, the exact pgaudit configuration is left to the user to define. pgaudit allows you to configure auditing rules either in `postgresql.conf` or within your SQL script.

For this test, we place pgaudit statements within a SQL script and verify that auditing is enabled and working. If you choose to configure pgaudit via a `postgresql.conf` file, then you will need to define your own custom `postgresql.conf` file and mount it to override the default `postgresql.conf` file.

To run this example, run the following to create a database container:

```
cd ./examples/docker/pgaudit
./run.sh
```

This starts a database on port 12005 on localhost. You can then run the test script as follows:

```
./test-pgaudit.sh
```

This test executes a SQL file which contains pgaudit configuration statements as well as executes some basic SQL commands. These SQL commands will cause pgaudit to create audit log messages in the `pg_log` log file created by the database container.

swarm - docker swarm

This example shows how to run a master and replica database container on a Docker Swarm (v.1.12) cluster.

First, set up a cluster, the Kubernetes libvirt coreos cluster example works well, see [coreos-libvirt-cluster](#)

Next, on each node, create the Swarm using these [Swarm Install](#)

Includes the command on the manager node:

```
docker swarm init --advertise-addr 192.168.10.1
```

Then the command on all the worker nodes:

```
docker swarm join \
  --token SWMTKN-1-65cn5wa1qv76l8l45uvlsbprogyhlprjpn27p1qxjwqmncn37o-
  015egopg4jhtbmlu04faon82u \
  192.168.10.1.37
```

Before creating Swarm services, for service discovery you need to define an overlay network to be used by the services you will create. Create the network like this:

```
docker network create --driver overlay crunchynet
```

We want to have the master database always placed on a specific node, we accomplish this using node constraints as follows:

```
docker node inspect kubernetes-node-1 | grep ID
docker node update --label-add type=master 18yrb7m650umx738rtevojpqy
```

In the above example, the kubernetes-node-1 node with ID 18yrb7m650umx738rtevojpqy has a user defined label of **master** added to it. The master service specifies **master** as a constraint when created, this tells Swarm to place the service on that specific node. The replica specifies a constraint of **node.labels.type != master** to have the replica always placed on a node that is not hosting the master service.

After you set up the Swarm cluster, you can then run the **examples/docker/swarm-service** example as follows on the **Swarm Manager Node**:

```
cd examples/docker/swarm-service
./run.sh
```

You can then find the nodes that are running the master and replica containers by:

```
docker service ps master
docker service ps replica
```

Given the PostgreSQL replica service is named **replica**, you can scale up the number of replica containers by running this command:

```
docker service scale replica=2
docker service ls
```

You can verify you have two replicas within PostgreSQL by viewing the **pg_stat_replication** table, the password is **password**, when logged into the kubernetes-node-1 host:

```
docker exec -it $(docker ps -q) psql -U postgres -c 'table pg_stat_replication'
postgres
```

You should see a row for each replica along with its replication status.

watch - automated failover

This example shows how to run the crunchy-watch container to perform an automated failover. For the example to work, the host on which you are running needs to allow read-write access to `/run/docker.sock`. The crunchy-watch container runs as the postgres user, so adjust the file permissions of `/run/docker.sock` accordingly.

Run the example as follows, (depends on master-replica example being run prior):

```
cd examples/docker/watch
./run.sh
```

This will start the watch container which tests every few seconds whether the master database is running, if not, it will trigger a failover (using docker exec) on the replica host.

Test it out by stopping the master:

```
docker stop master
docker logs watch
```

Look at the watch container logs to see it perform the failover.

Tips - send a signal to postgres

first, find the PID of the postmaster:

```
docker exec -it master cat /pgdata/master/postmaster.pid
```

then, send it the signal to kill it or other signal depending on what you want to do:

```
docker exec -it master kill -SIGTERM 22
```

Legal Notices

Copyright © 2017 Crunchy Data Solutions, Inc.

CRUNCHY DATA SOLUTIONS, INC. PROVIDES THIS GUIDE "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF NON INFRINGEMENT, MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Crunchy, Crunchy Data Solutions, Inc. and the Crunchy Hippo Logo are trademarks of Crunchy Data Solutions, Inc.