Openshift Examples - Crunchy Containers

Crunchy Data Solutions, Inc.

Version 1.2.8, 2017-02-10

Examples for the Openshift Environment

The examples/openshift directory containers examples for running the Crunchy containers in an Openshift environment.

The examples are explained below.

Some of the examples make use of NFS for a persistent volume file system type, see the NFS Note at the bottom of this document for details on how NFS is configured in these examples.

Notes on NFS installation and file system permissions are in the install.asciidoc documentation.

For running the examples that require persistent volumes, you will need to run the following script:

```
cd $BUILDBASE/examples/openshift
./create-pv.sh
```

basic - Running a single database container

This openshift template will create a single master PostgreSQL instance.

Running the example:

```
cd examples/openshift/basic
./run.sh
```

You can see what passwords were generated by running this command:

```
oc describe pod basic | grep PG
```

You can run the following command to test the database, entering the value of PG_PASSWORD from above for the password when prompted:

```
psql -h basic.openshift.svc.cluster.local -U testuser userdb
```

pgpool - pgpool for master replica examples

You can create a pgpool service that will work with the master and replica created in the previous Example 1 and Example 2.

You will need to edit the pgpool-for-master-replica-rc-dc-replicas-only.json and supply the testuser password that was generated when you created the master replica pods, then run the following command to deploy the pgpool service:

```
cd examples/openshift/pgpool
./run.sh
```

Next, you can access the master replica cluster via the pgpool service by entering the following command:

```
psql -h pgpool-rc -U testuser userdb
psql -h pgpool-rc -U testuser postgres
```

when prompted, enter the password for the PG_USERNAME testuser that was set for the pg-master pod, typically it is **password**.

At this point, you can enter SELECT and INSERT statements and pgpool will proxy the SQL commands to the master or replica(s) depending on the type of SQL command. Writes will always be sent to the master, and reads will be sent (round-robin) to the replica(s).

You can view the nodes that pgpool is configured for by running:

```
psql -h pgpool-rc -U testuser userdb -c 'show pool_nodes'
```

master-replica-rc-dc - master replica with scalable replicas

This example is similar to the previous examples but builds a master pod, and a single replica that can be scaled up using a replication controller. The master is implemented as a single pod since it can not be scaled like read-only replicas.

Running the example:

```
oc create -f master-replica-rc-dc-replicas-only.json | oc create -f -
```

Connect to the PostgreSQL instances with the following:

```
psql -h pg-master-rc-dc.pgproject.svc.cluster.local -U testuser userdb
psql -h pg-replica-rc-dc.pgproject.svc.cluster.local -U testuser userdb
```

Here is an example of increasing or scaling up the Postgres 'replica' pods to 2:

Here is an example of increasing or scaling up the Postgres 'replica' pods to 2:

Here is an example of increasing or scaling up the Postgres 'replica' pods to 2:

```
oc scale rc pg-replica-rc-dc-1 --replicas=2
```

Enter the following commands to verify the PostgreSQL replication is working.

```
psql -c 'table pg_stat_replication' -h pg-master-rc.pgproject.svc.cluster.local -U
master postgres
psql -h pg-replica-rc.pgproject.svc.cluster.local -U master postgres
```

You can see that the replica service is load balancing between multiple replicas by running a command as follows, run the command multiple times and the ip address should alternate between the replicas:

```
psql -c 'select inet_server_addr()' -h pg-replica-rc -U master postgres
```

backup - Performing a Full Backup

This example assumes you have a database pod running called **basic** as created by the **basic** example and that you have configured NFS as decribed in the Prerequisites section of this document..

You can perform a database backup by executing the following step:

```
cd examples/openshift/backup-job
./run.sh
```

A successful backup will perform pg_basebackup on the pg-master and store the backup in the NFS mounted volume under a directory named pg-master, each backup will be stored in a subdirectory with a timestamp as the name. This allows any number of backups to be kept.

The **examples/openshift/crunchy-pv-backup.json** specifies a **persistentVolumeReclaimPolicy** of **Retain** to tell Openshift that we want to keep the volume contents after the removal of the PV.

master-nfs - NFS Example

This example will create a single master postgres pod that is using an NFS volume to store the postgres data files.

master-restore - example of restoring a database from a backup

This is an example of restoring a database pod using an existing backup archive located on an NFS volume.

First, locate the database backup you want to restore, for example:

```
/nfsfileshare/pg-master/2016-01-29:22:34:20
```

Then create the pod:

```
./run.sh
```

When the database pod starts, it will copy the backup files to the database directory inside the pod and start up postgres as usual.

The restore only takes place if:

- the /pgdata directory is empty
- the /backups directory contains a valid postgresql.conf file

Openshift Example 7 - Failover Example

An example of performing a database failover is described in the following steps:

- create a master and replica replication using master-replica-rc-dc-replicas-only.json

```
oc process -f master-replica-rc-dc-replicas-only.json | oc create -f -
```

- scale up the number of replicas to 2

```
oc scale rc pg-replica-rc-1 --replicas=2
```

- delete the master pod

```
oc delete pod pg-master-rc
```

- exec into a replica and create a trigger file to being the recovery process, effectively turning the replica into a master

```
oc exec -it pg-replica-rc-1-lt5a5
touch /tmp/pg-failover-trigger
```

- change the label on the replica to pg-master-rc instead of pg-replica-rc

```
oc edit pod/pg-replica-rc-1-lt5a5
original line: labels/name: pg-replica-rc
updated line: labels/name: pg-master-rc
```

```
or alternatively:
```

```
oc label --overwrite=true pod pg-replica-rc-1-lt5a5 name=pg-master-rc
```

You can test the failover by creating some data on the master and then test to see if the replicas have the replicated data.

```
psql -c 'create table foo (id int)' -U master -h pg-master-rc postgres
psql -c 'table foo' -U master -h pg-replica-rc postgres
```

After a failover, you would most likely want to create a database backup and be prepared to recreate your cluster from that backup.

master-replica-rc-nfs - Master Slave Deployment using NFS

This example uses NFS volumes for the master and the replicas. In some scenarios, customers might want to have all the Postgres instances using NFS volumes for persistence.

Relevant files for this example:

- master-replica-rc-nfs.json This file creates the master and replica deployment, creating pods and services where the replica is controlled by a Replication Controller, allowing you to scale up the replicas.

To run the example, follow these steps:

As the project user, create the master replica deployment:

```
./run.sh
```

If you examing your NFS directory, you will see postgres data directories created and used by your master and replica pods.

Next, add some test data to the master:

```
psql -c 'create table testtable (id int)' -U master -h pg-master-rc-nfs postgres
psql -c 'insert into testtable values (123)' -U master -h pg-master-rc-nfs postgres
```

Next, add a new replica:

```
oc scale rc pg-replica-rc-nfs-1 --replicas=2
```

At this point, you should see the new NFS directory created by the new replica pod, and you should also be able to test that replication is working on the new replica:

```
psql -c 'table testtable' -U master -h pg-replica-rc-nfs postgres
```

badger - pgbadger example

This example creates a pod that contains a database container and a pgbadger container.

pgbadger is then served up on port 10000. Each time you do a GET on http://pg-master:10000/api/badgergenerate it will run pgbadger against the database log files running in the pg-master container.

To run the example:

```
cd examples/openshift/badger
./run.sh
```

try the following command to see the generated HTML output:

```
curl http://badger-example:10000/api/badgergenerate
```

You can view this output in a browser if you allow port forwarding from your container to your server host using a command like this:

```
socat tcp-listen:10001,reuseaddr,fork tcp:pg-master:10000
```

This command maps port 10000 of the service/container to port 10001 of the local server. You can now use your browser to see the badger report.

This is a short-cut way to expose a service to the external world, Openshift would normally configure a Router whereby you could 'expose' the service in an Openshift way. Here is the docs on installing the Openshift Router:

```
https://docs.openshift.com/enterprise/3.0/install_config/install/deploy_router.html
```

secret - database with secrets

This example allows you to set the Postgresql passwords using Kube Secrets.

The secret uses a base64 encoded string to represent the values to be read by the container during initialization. The encoded password value is **password**. Run the example as follows:

```
examples/openshift/secret/run.sh
```

The secrets are mounted in the **/pguser**, **/pgmaster**, **/pgroot** volumes within the container and read during initialization. The container scripts create a Postgres user with those values, and sets the passwords for the master user and postgres superuser using the mounted secret volumes.

When using secrets, you do NOT have to specify the following env vars if you specify all three secrets volumes: * PG_USER * PG_PASSWORD * PG_ROOT_PASSWORD * PG_MASTER_USER * PG_MASTER_PASSWORD

You can test the container as follows, in all cases, the password is **password**:

```
psql -h secret-pg -U pguser1 postgres
psql -h secret-pg -U postgres postgres
psql -h secret-pg -U master postgres
```

watch - Automated Failover

This example shows how a form of automated failover can be configured for a master and replica deployment.

First, create a master and a replica, in this case the replica lives in a Deployment which can scale up:

```
cd examples/openshift/master-replica-dc
./run.sh
```

Next, create an Openshift service account which is used by the crunchy-watch container to perform the failover, also set policies that allow the service account the ability to edit resources within the openshift and default projects :

```
oc create -f sa.json
oc policy add-role-to-group edit system:serviceaccounts -n openshift
oc policy add-role-to-group edit system:serviceaccounts -n default
```

Next, create the container that will 'watch' the Postgresql cluster:

```
cd examples/openshift/watch
./run.sh
```

At this point, the watcher will sleep every 20 seconds (configurable) to see if the master is responding. If the master doesn't respond, the watcher will perform the following logic:

- log into openshift using the service account

- set its current project

- find the first replica pod

- delete the master service saving off the master service definition

- create the trigger file on the first replica pod

- wait 20 seconds for the failover to complete on the replica pod

- edit the replica pod's lable to match that of the master

- recreate the master service using the stored service definition

- loop through the other remaining replica and delete its pod

At this point, clients when access the master's service will actually be accessing the new master. Also, Openshift will recreate the number of replicas to its original configuration which each replica pointed to the new master. Replication from the master to the new replicas will be started as each new replica is started by Openshift.

To test it out, delete the master pod and view the watch pod log:

```
oc delete pod pg-master-rc-dc
oc logs watch
oc get pod
```

metrics - Metrics Collection

This example shows how postgres metrics can be collected and stored in prometheus and graphed with grafana.

First, create the crunchy-metrics pod which contains the prometheus data store and the grafana graphing web application:

```
./run.sh
```

At this point, you can view the prometheus web console at crunchy-metrics:9090, the prometheus push gateway at crunchy-metrics:9091, and the grafana web app at crunchy-metrics:3000.

Next, start a postgres pod that has the crunchy-collect container as follows:

```
cd $BUILDBASE/examples/openshift/collect
./run.sh
```

At this point, metrics will be collected every 3 minutes and pushed to prometheus. You can build graphs off the metrics using grafana.

vacuum - Vacuum job

This example shows how you can run a vacuum job against a postgres database container.

The crunchy-vacuum container image exists to allow a DBA a way to run a job either one-off or scheduled to perform a variety of vacuum operations.

To run the vacuum a single time, an example is included as follows from the examples/openshift directory:

```
cd examples/openshift/master-replica
./run.sh
cd ../vacuum-job
./run.sh
```

This will start a vacuum container that runs as a Kube Job type. It will run once. The crunchy-vacuum image is executed, passed in the Postgres connection parameters to the single-master postgres container. The type of vacuum performed is dictated by the environment variables passed into the job. The complete set of environment variables read by the vacuum job include:

- VAC_FULL - when set to true adds the FULL parameter to the VACUUM command
- VAC_TABLE - when set, allows you to specify a single table to vacuum, when not specified, the entire database tables are vacuumed
- JOB_HOST - required variable is the postgres host we connect to
- PG_USER - required variable is the postgres user we connect with
- PG_DATABASE - required variable is the postgres database we connect to
- PG_PASSWORD - required variable is the postgres user password we connect with
- PG_PORT - allows you to override the default value of 5432
- VAC_ANALYZE - when set to true adds the ANALYZE parameter to the VACUUM command
- VAC_VERBOSE - when set to true adds the VERBOSE parameter to the VACUUM command
- VAC_FREEZE - when set to true adds the FREEZE parameter to the VACUUM command

custom-config - Custom Configuration Files

This example shows how you can use your own customized version of setup.sql when creating a postgres database container.

If you mount a /pgconf volume, crunchy-postgres will look at that directory for postgresql.conf, pg_hba.conf, and setup.sql. If it finds one of them it will use that file instead of the default files. Currently, if you specify a postgresql.conf file, you also need to specify a pg_hba.conf file.

The example shows how a custom setup.sql file can be used. Run it as follows from the examples/openshift/custom-config directory:

```
./run.sh
```

This will start a database container that will use an NFS mounted /pgconf directory that will container the custom setup.sql file found in the example directory.

custom-config-sync - Custom Configuration Files with Sync Replica

This example shows how you can use your own customized version of postgresql.conf and pg_hba.conf to override the default configuration. It also specifies a sync replica in the postgresql.conf and starts up a sync replica.

If you mount a /pgconf volume, crunchy-postgres will look at that directory for postgresql.conf, pg_hba.conf, and setup.sql. If it finds one of them it will use that file instead of the default files. Currently, if you specify a postgresql.conf file, you also need to specify a pg_hba.conf file.

Run it as follows from the examples/openshift/custom-config-sync directory:

```
./run.sh
```

This will start a **csmaster** container that will use the custom config files when the database is running. It will also create a sync replica named **cssyncreplica**, this replica is connected to the master via streaming replication.

pgbouncer - pgbouncer

This example shows how you can use the crunchy-pgbouncer container when running under Openshift.

The example assumes you have run the master/replica example found here:

```
examples/openshift/master-replica-dc
```

Then you would start up the pgbouncer container using the following example:

```
examples/openshift/pgbouncer
```

The example assumes you have an NFS share path of /nfsfileshare/! NFS is required to mount the pgbouncer configuration files which are then mounted to /pgconf in the crunchy-pgbouncer container.

If you mount a /pgconf volume, crunchy-postgres will look at that directory for postgresql.conf, pg_hba.conf, and setup.sql. If it finds one of them it will use that file instead of the default files.

Test the example by killing off the master database container as follows:

```
oc delete pod pg-master-rc-dc
```

Then watch the pgbouncer log as follows to confirm it detects the loss of the master:

```
oc logs pgbouncer
```

After the failover is completed, you should be able to access the new master using the master service as follows:

```
psql -h pg-master-rc-dc.openshift.svc.cluster.local -U master postgres
```

and access the replica as follows:

```
psql -h pg-replica-rc-dc.openshift.svc.cluster.local -U master postgres
```

or via the pgbouncer proxy as follows:

```
psql -h pgbouncer.openshift.svc.cluster.local  -U master master
```

sync - synchrounous replica

This example deploys a PostgreSQL cluster with a master, a synchrounous replica, and an asynchronous replica. The two replicas share the same Service.

Running the example:

```
examples/openshift/sync/run.sh
```

Connect to the **master** and **replica** databases as follows:

```
psql -h master -U postgres postgres -c 'create table mister (id int)'
psql -h master -U postgres postgres -c 'insert into mister values (1)'
psql -h master -U postgres postgres -c 'table pg_stat_replication'
psql -h replica -U postgres postgres -c 'select inet_server_addr(), * from mister'
psql -h replica -U postgres postgres -c 'select inet_server_addr(), * from mister'
psql -h replica -U postgres postgres -c 'select inet_server_addr(), * from mister'
```

This set of queries will show you the IP address of the Postgres replica container, notice it changes because of the round-robin Service proxy we are using for both replicas. The example queries also show that both replicas are replicating from the master.

pgadmin4 - pgadmin4

This example, examples/openshift/pgadmin4, provides a container that runs the pgadmin4 web application.

To run this example, run the following:

```
cd $BUILDBASE/examples/openshift/pgadmin4
./run.sh
```

This script creates the **pgadmin4** pod and service, it will expose port 5050.

You should now be able to browse to http://pgadmin4.openshift.svc.cluster.local:5050 and log into the web application using a user ID of **admin@admin.org** and password of **password**. Replace YOURLOCALIP with whatever your local IP address happens to be.

workshop - workshop

This example, examples/openshift/workshop, provides an example of using Openshift Templates to build pods, routes, services, etc.

You use the **oc new-app** command to create objects from the JSON templates. This is an alternative way to create Openshift objects instead of using **oc create**.

This example is used within a joint Redhat-Crunchy workshop that is given at various conferences to demonstrate Openshift and Crunchy Containers working together. Thanks to Steven Pousty from Redhat for this example!

See the README file within the workshop directory for instructions on running the example.

pitr - PITR (point in time recovery)

This is a complex example. For details on how PITR is implemented within the Suite, see the PITR Documentation for details and background.

This example, examples/openshift/pitr, provides an example of performing a PITR using Openshift.

Lets start by running the example database container:

```
cd $BUILDBASE/examples/openshift/pitr
./run-master-pitr.sh
```

This step will create a database container, **master-pitr**. This container is configured to continuously write WAL segment files to a mounted volume (/pgwal).

After you start the database, you will create a base backup using this command:

```
./run-master-pitr-backup.sh
```

This will create a backup and write the backup files to a persistent volume (/pgbackup).

Next, lets create some recovery targets within the database, run the SQL commands against the **master-pitr** database as follows:

```
./run-sql.sh
```

This will create recovery targets named **beforechanges**, **afterchanges**, and **nomorechanges**. It will create a table, **pitrtest**, between the **beforechanges** and **afterchanges** targets. It will also run a SQL CHECKPOINT to flush out the changes to WAL segments.

Next, now that we have a base backup and a set of WAL files containing our database changes, we can shut down the **master-pitr** database to simulate a database failure. Do this by running the following:

```
oc delete pod master-pitr
```

Next, we will create 3 different restored database containers based upon the base backup and the saved WAL files.

First, we restore prior to the **beforechanges** recovery target. This recovery point is **before** the **pitrtest** table is created.

Edit the master-pitr-restore.json file, and edit the environment variable to indicate we want to use the **beforechanges** recovery point:

```
}, {
"name": "RECOVERY_TARGET_NAME",
"value": "beforechanges"
}, {
```

Then run the following to create the restored database container:

```
./run-restore-pitr.sh
```

After the database has restored, you should be able to perform a test to see if the recovery worked as expected:

```
psql -h master-pitr-restore.openshift.svc.cluster.local -U postgres postgres -c 'table
pitrtest'
psql -h master-pitr-restore.openshift.svc.cluster.local -U postgres postgres -c
'create table foo (id int)'
psql -h master-pitr-restore.openshift.svc.cluster.local -U postgres postgres -c
'select pg_xlog_replay_resume()'
psql -h master-pitr-restore.openshift.svc.cluster.local -U postgres postgres -c
'create table foo (id int)'
```

The output of these command should show that the **pitrtest** table is not present. It should also show that you can not create a new table because the database is paused in recovery mode. Lastly, if you execute a **resume** command, it will show that you can now create a table as the database has fully recovered.

You can also test that if **afterchanges** is specified, that the **pitrtest** table is present but that the database is still in recovery mode.

Lastly, you can test a full recovery using **all** of the WAL files, if you remove the **RECOVERY_TARGET_NAME** environment variable completely.

The NFS portions of this example depend upon an NFS file system with the following path configurations be present:

```
/nfsfileshare
/nfsfileshare/backups
/nfsfileshare/WAL
```

backrest - pgbackrest example

This example shows how to enable pgbackrest as the archiver within the crunchy-postgres container. See the pgbackrest Documentation for details and background.

Start by running the example database container:

```
cd examples/openshift/backrest
./run.sh
```

This will create the following: * PV/PVC for /pgconf and /backrestrepo volumes * master database pod * master service

The run.sh script copies the pgbackrest.conf configuration file to /nfsfileshare/pgconf which is our NFS file path.

The archive files are written to the NFS path of /nfsfileshare/backrestrepo.

The presence of /pgconf/pgbackrest.conf is what is used to determine whether pgbackrest will be used as the archive command or not. You will need to specify the ARCHIVE_TIMEOUT environment variable as well to use this.

After you run the example, you should see archive files being written to the /backrestrepo volume (/nfsfileshare/backrestrepo).

configmap- database credentials from a configmap

This example shows how to use a configmap to store the postgresql.conf and pg_hba.conf files to be used when overriding the default configuration within the container.

Start by running the database container:

```
cd $BUILDBASE/examples/openshift/configmap
./run.sh
```

The files, pg_hba.conf and postgresql.conf, in the example directory are used to create a configmap object within OpenShift. Within the run.sh script, the configmap is created, and notice within the configmap.json file how the /pgconf mount is related to the configmap.

Openshift Tips

Tip : Finding the Postgresql Passwords

The passwords used for the PostgreSQL user accounts are generated by the Openshift 'process' command. To inspect what value was supplied, you can inspect the master pod as follows:

```
oc get pod pg-master-rc-1-n5z8r -o json
```

Look for the values of the environment variables: - PG_USER - PG_PASSWORD - PG_DATABASE

Tip : Examining a backup job log

Database backups are implemented as a Kubernetes Job. A Job is meant to run one time only and not be restarted by Kubernetes. To view jobs in Openshift you enter:

```
oc get jobs
oc describe job backupjob
```

You can get detailed logs by referring to the pod identifier in the job 'describe' output as follows:

```
oc logs backupjob-pxh2o
```

Tip : Password Mgmt

Remember that if you do a database restore, you will get whatever user IDs and passwords that were saved in the backup. So, if you do a restore to a new database and use generated passwords, the new passwords will not be the same as the passwords stored in the backup!

You have various options to deal with managing your passwords.

- externalize your passwords using secrets instead of using generated values
- manually update your passwords to your known values after a restore

Note that you can edit the environment variables when there is a 'dc' using, currently only the replicas have a 'dc' to avoid the possiblity of creating multiple masters, this might need to change in the future, to better support password management:

```
oc env dc/pg-master-rc PG_MASTER_PASSWORD=foo PG_MASTER=user1
```

Tip : Log Aggregation

Openshift can be configured to include the EFK stack for log aggregation. Openshift Administrators can configure the EFK stack as documented here:

https://docs.openshift.com/enterprise/3.1/install_config/aggregate_logging.html

Tip : nss_wrapper

If an Openshift deployment requires that random generated UIDs be supported by containers, the Crunchy containers can be modifed similar to those located here to support the use of nss_wrapper to equate the random generated UIDs/GIDs by openshift with the postgres user:

https://github.com/openshift/postgresql/blob/master/9.4/root/usr/share/container-scripts/postgresql/common.sh

Tip : encoding secrets

You can use kubernetes secrets to set and maintain your database credentials. Secrets requires you base64 encode your user and password values as follows:

```
echo -n 'myuserid' | base64
```

You will paste these values into your JSON secrets files for values.

docker to be installed.

You can keep yum from upgrading docker by including this line in your /etc/yum.conf file:

```
exclude=docker-1.9* docker-selinux-1.9*
```

Tip : DNS configuration for Openshift development

As of OSE 3.3, the following DNS modifications are not typically necessary any longer....but I'm leaving them here as a reference....

Luke Meyer from Redhat wrote an excellent blog on how to configure dnsmasq and Openshift, it is located here:

http://developers.redhat.com/blog/2015/11/19/dns-your-openshift-v3-cluster/

Key things included in this blog are:

- configuring dhcp to include the local IP address in /etc/resolv.conf upon boot
- configuring dnsmasq
- configuring openshift dns to listen on another port

In my dev setup, I have openshifts DNS listening on 127.0.0.1:8053. I have my dnsmasq listening on the local IP address 192.168.0.109:53

Therefore in my /etc/dhcp/dhclient.conf I have this config:

```
prepend domain-name-servers 192.168.0.109;
```

If you dont have your DNS configured correctly, replication controllers and deployment configs basically will not work.

Tip : system policys for pv creation and listing

For my testing, I wanted to allow the **system** user to be able to create and list persistent volumes, as of OSE 3.3, I had to enter these commands as the **root** user after installation to modify the policies:

```
oadm policy add-role-to-user cluster-reader system
oc describe clusterPolicyBindings :default
oadm policy add-cluster-role-to-user cluster-reader system
oc describe clusterPolicyBindings :default
oc describe clusterPolicyBindings :default
oadm policy add-cluster-role-to-user cluster-admin system
```

Tip : anyuid permissions

For my testing, I create a user named **test** on OSE, then I run the following command to grant it permission to use the **anyuid** SCC:

```
oc adm policy add-scc-to-group anyuid system:authenticated
```

This says that any authenticate user can run with the anyuid SCC which lets them create PVCs and use the **fsGroup** setting for the Postgres containers to work using NFS. There is probably a smarter and more precise way to grant this permission?

Tip 15: NFS Setup

To control the permissions of the NFS file system certain examples make use of the **supplementalGroups** security context setting for pods. In these examples, we specify the GID of the **nfsnobody** group (65534). If you want to use a different GID for the supplementalGroup then you will need to alter the NFS examples accordingly.

When the pod runs, the pod user is UID **26** which is the postgres user ID. By specifying the **supplementalGroup** the pod will also be added to the **nfsnobody** group. So, when you set up your NFS mount, you can specify the permissions to be as follows:

```
drwxrwx---.   3 nfsnobody nfsnobody   23 Dec 16 11:28 nfsfileshare
```

This restricts **other** users from writing to the NFS share, but will allow the **nfsnobody** group to have write access. This way, the NFS mount permissions can be managed to only allow certain pods write access.

Also, remember that on systems with SELinux set to enforcing mode that you will need to allow NFS write permissions by running this command:

```
sudo setsebool -P virt_use_nfs 1
```

Note that supplementalGroup settings are required for NFS but you would use the fsGroup setting for the AWS file system. Check out this link for details: https://docs.openshift.org/latest/install_config/persistent_storage/pod_security_context.html.

Legal Notices