



Openshift Examples - Crunchy Containers

Crunchy Data Solutions, Inc.

Version 1.2.3, 2016-09-27

Examples for the Openshift Environment

The examples/openshift directory contains examples for running the Crunchy containers in an Openshift environment.

The examples are explained below.

Openshift Example 1 - simple master

This openshift template will create a single master PostgreSQL instance.

Running the example:

```
cd single-master  
./run.sh
```

You can see what passwords were generated by running this command:

```
oc describe pod single-master | grep PG
```

You can run the following command to test the database, entering the value of PG_PASSWORD from above for the password when prompted:

```
psql -h single-master.openshift.svc.cluster.local -U testuser userdb
```

Openshift Example 2 - slave.json

This Openshift template will create a single slave PostgreSQL instance that will connect to the master created in Example 1.

You will need to edit the slave.json file and enter the PG_MASTER_PASSWORD value generated in Example 1 as found in the master environment variables.

Running the example:

```
oc create -f slave.json | oc create -f -
```

You can run the following command to test the database, entering the value of PG_PASSWORD from above for the password when prompted:

```
psql -c 'select * from pg_stat_replication' -h pg-master.pgproject.svc.cluster.local -U master userdb
psql -c 'create table foo (id int)' -h pg-master.pgproject.svc.cluster.local -U master userdb
psql -c 'insert into foo values (123)' -h pg-master.pgproject.svc.cluster.local -U master userdb
psql -c 'table foo' -h pg-slave.pgproject.svc.cluster.local -U master userdb
```

If replication is working, you should see a row returned in the `pg_stat_replication` table query and also a row returned from the `pg-slave` query.

Openshift Example 3 - pgpool for master slave examples

You can create a `pgpool` service that will work with the master and slave created in the previous Example 1 and Example 2.

You will need to edit the `pgpool-for-master-slave-rc-dc-slaves-only.json` and supply the `testuser` password that was generated when you created the master slave pods, then run the following command to deploy the `pgpool` service:

```
cd examples/openshift/pgpooltest
./run.sh
```

Next, you can access the master slave cluster via the `pgpool` service by entering the following command:

```
psql -h pgpool-rc -U testuser userdb
psql -h pgpool-rc -U testuser postgres
```

when prompted, enter the password for the `PG_USERNAME testuser` that was set for the `pg-master` pod, typically it is **password**.

At this point, you can enter `SELECT` and `INSERT` statements and `pgpool` will proxy the SQL commands to the master or slave(s) depending on the type of SQL command. Writes will always be sent to the master, and reads will be sent (round-robin) to the slave(s).

You can view the nodes that `pgpool` is configured for by running:

```
psql -h pgpool-rc -U testuser userdb -c 'show pool_nodes'
```

Openshift Example 4 - master-slave-rc-dc-slaves-only.json

This example is similar to the previous examples but builds a master pod, and a single slave that can be scaled up using a replication controller. The master is implemented as a single pod since it can not be scaled like read-only slaves.

Running the example:

```
oc create -f master-slave-rc-dc-slaves-only.json | oc create -f -
```

Connect to the PostgreSQL instances with the following:

```
psql -h pg-master-rc-dc.pgproject.svc.cluster.local -U testuser userdb  
psql -h pg-slave-rc-dc.pgproject.svc.cluster.local -U testuser userdb
```

Here is an example of increasing or scaling up the Postgres 'slave' pods to 2:

Here is an example of increasing or scaling up the Postgres 'slave' pods to 2:

Here is an example of increasing or scaling up the Postgres 'slave' pods to 2:

```
oc scale rc pg-slave-rc-dc-1 --replicas=2
```

Enter the following commands to verify the PostgreSQL replication is working.

```
psql -c 'table pg_stat_replication' -h pg-master-rc.pgproject.svc.cluster.local -U  
master postgres  
psql -h pg-slave-rc.pgproject.svc.cluster.local -U master postgres
```

You can see that the slave service is load balancing between multiple slaves by running a command as follows, run the command multiple times and the ip address should alternate between the slaves:

```
psql -c 'select inet_server_addr()' -h pg-slave-rc -U master postgres
```

Openshift Example 5 - Performing a Backup

This example assumes you have a master database pod running called pg-master as created by Openshift Example 1 and that you have configured NFS as described in the Prerequisites section of this document..

You can perform a database backup by executing the following step:

- as openshift admin, run:

```
oc create -f pv-backups.json
```

- as normal user, run:

```
oc create -f pvc-backups.json  
oc create -f backup-job-nfs.json
```

A successful backup will perform `pg_basebackup` on the `pg-master` and store the backup in the NFS mounted volume under a directory named `pg-master`, each backup will be stored in a subdirectory with a timestamp as the name. This allows any number of backups to be kept.

The `pv-backups.json` specifies a `ReclaimPolicy` of `Retain` to tell Openshift that we want to keep the volume contents after the removal of the PV.

Lastly, in order to rerun this same backup, you will need to remove the Job, PVC, and PV used in this example, this is done as follows:

- as admin user

```
oc delete pv backup-pg-master
```

- as normal user

```
oc delete pvc backup-claim-pg-master  
oc delete job backupjob-pg-master
```

Openshift Example 5 - NFS Example

I have provided an example of using NFS for the postgres data volume. On my test nfs server, I had to set the exports file entry as follows:

```
/nfsfileshare * (rw,insecure,sync)
```

First, you can only create persistent volumes as a cluster admin, you can login in as the admin user as follows:

```
oc login -u system:admin
```

To run it, you would execute the following as the openshift administrator:

```
oc create -f master-nfs-pv.json
```

Then as the normal openshift user account, create the Persistence Volume Claim and database pod as follows:

```
oc create -f master-nfs-pvc.json  
oc process -f master-nfs.json | oc create -f -
```

This will create a single master postgres pod that is using an NFS volume to store the postgres data files.

Openshift Example 6 - Restore Example

I have provided an example of restoring a database pod using an existing backup archive located on an NFS volume.

First, locate the database backup you want to restore, for example:

```
/nfsfileshare/pg-master/2016-01-29:22:34:20
```

Next, * edit the master-restore-pv.json file to use that path in building the PV, * edit the master-restore-pv.json file to use a unique label * and then execute as the openshift superuser:

```
oc login -u system:admin  
oc create -f master-restore-pv.json
```

Next, * edit the master-restore-pvc.json file, specify the same unique label used in the master-restore-pv.json file. * Then execute as the normal test user:

```
oc create -f master-restore-pvc.json
```

Next, create a database pod as the normal user:

```
oc process -f master-restore.json | oc create -f -
```

When the database pod starts, it will copy the backup files to the database directory inside the pod and start up postgres as usual.

The restore only takes place if:

- the /pgdata directory is empty
- the /backups directory contains a valid postgresql.conf file

Openshift Example 7 - Failover Example

An example of performing a database failover is described in the following steps:

- create a master and slave replication using master-slave-rc-dc-slaves-only.json

```
oc process -f master-slave-rc-dc-slaves-only.json | oc create -f -
```

- scale up the number of slaves to 2

```
oc scale rc pg-slave-rc-1 --replicas=2
```

- delete the master pod

```
oc delete pod pg-master-rc
```

- exec into a slave and create a trigger file to begin the recovery process, effectively turning the slave into a master

```
oc exec -it pg-slave-rc-1-lt5a5  
touch /tmp/pg-failover-trigger
```

- change the label on the slave to pg-master-rc instead of pg-slave-rc

```
oc edit pod/pg-slave-rc-1-lt5a5  
original line: labels/name: pg-slave-rc  
updated line: labels/name: pg-master-rc
```

or alternatively:

```
oc label --overwrite=true pod pg-slave-rc-1-lt5a5 name=pg-master-rc
```

You can test the failover by creating some data on the master and then test to see if the slaves have the replicated data.

```
psql -c 'create table foo (id int)' -U master -h pg-master-rc postgres  
psql -c 'table foo' -U master -h pg-slave-rc postgres
```

After a failover, you would most likely want to create a database backup and be prepared to recreate your cluster from that backup.

Openshift Example 8 - Master Slave Deployment using NFS

This example uses NFS volumes for the master and the slaves. In some scenarios, customers might want to have all the Postgres instances using NFS volumes for persistence.

Relevant files for this example:

- master-slave-rc-nfs.json This file creates the master and slave deployment, creating pods and services where the slave is controlled by a Replication Controller, allowing you to scale up the slaves.

To run the example, follow these steps:

- as the openshift admin, create the required PV(s) using this command:

```
oc create -f master-slave-rc-nfs-pv.json
oc create -f master-slave-rc-nfs-pv2.json
```

This will create a PV for the master and another PV for the slaves. * as the project user, create the required PVC(s) using this command:

```
oc create -f master-slave-rc-nfs-pvc2.json
oc create -f master-slave-rc-nfs-pvc.json
```

This will create a PVC for the master and another PVC for the slaves. * as the project user, create the master slave deployment:

```
oc process -f master-slave-rc-nfs.json | oc create -f -
```

If you examine your NFS directory, you will see postgres data directories created and used by your master and slave pods.

Next, add some test data to the master:

```
psql -c 'create table testtable (id int)' -U master -h pg-master-rc-nfs postgres
psql -c 'insert into testtable values (123)' -U master -h pg-master-rc-nfs postgres
```

Next, add a new slave:

```
oc scale rc pg-slave-rc-nfs-1 --replicas=2
```

At this point, you should see the new NFS directory created by the new slave pod, and you should also be able to test that replication is working on the new slave:


```
psql -c 'table testtable' -U master -h pg-slave-rc-nfs postgres
```

Openshift Example 9 - Master with pgbadger add-in

This example uses a version of master.json but also adds the pgbadger container to the pg-master pod. pgbadger is then served up on port 10000. Each time you do a GET on <http://pg-master:10000/api/badgergenerate> it will run pgbadger against the database log files running in the pg-master container.

To run the example:

```
cd examples/openshift/badger
./run.sh
```

try the following command to see the generated HTML output:

```
curl http://pg-master:10000/api/badgergenerate
```

You can view this output in a browser if you allow port forwarding from your container to your server host using a command like this:

```
socat tcp-listen:10001,reuseaddr,fork tcp:pg-master:10000
```

This command maps port 10000 of the service/container to port 10001 of the local server. You can now use your browser to see the badger report.

This is a short-cut way to expose a service to the external world, Openshift would normally configure a Router whereby you could 'expose' the service in an Openshift way. Here is the docs on installing the Openshift Router:

```
https://docs.openshift.com/enterprise/3.0/install\_config/install/deploy\_router.html
```

Openshift Example 10 - Master with readiness probe

This example uses a version of master.json but also adds a Kubernetes readiness probe specific for postgresql. This readiness probe uses the postgres pg_isready utility to attempt a connection to postgres from within the container using the postgres user and postgres database as the parameters.

Run the following:

```
oc create -f master-ready.json | oc create -f -
```

Openshift Example 11 - Master with secrets

This example allows you to set the Postgresql passwords using Kube Secrets.

The secret uses a base64 encoded string to represent the values to be read by the container during initialization. The encoded password value is **password**. Run the example as follows:

```
examples/openshift/secret/run.sh
```

The secrets are mounted in the **/pguser**, **/pgmaster**, **/pgroot** volumes within the container and read during initialization. The container scripts create a Postgres user with those values, and sets the passwords for the master user and postgres superuser using the mounted secret volumes.

When using secrets, you do NOT have to specify the following env vars if you specify all three secrets volumes: * PG_USER * PG_PASSWORD * PG_ROOT_PASSWORD * PG_MASTER_USER * PG_MASTER_PASSWORD

You can test the container as follows, in all cases, the password is **password**:

```
psql -h secret-pg -U pguser1 postgres
psql -h secret-pg -U postgres postgres
psql -h secret-pg -U master postgres
```

Openshift Example 12 - Automated Failover

This example shows how a form of automated failover can be configured for a master and slave deployment.

First, create a master and 2 slaves:

```
oc process -f master-slave.json
oc scale rc pg-slave-rc-dc-1 --replicas=2
```

Next, create an Openshift service account which is used by the crunchy-watch container to perform the failover, also set policies that allow the service account the ability to edit resources within the openshift and default projects :

```
oc create -f sa.json
oc policy add-role-to-group edit system:serviceaccounts -n openshift
oc policy add-role-to-group edit system:serviceaccounts -n default
```

Next, create the container that will 'watch' the Postgresql cluster:

```
oc process -f watch.json | oc create -f -
```

At this point, the watcher will sleep every 20 seconds (configurable) to see if the master is responding. If the master doesn't respond, the watcher will perform the following logic:

- log into openshift using the service account
- set its current project
- find the first slave pod
- delete the master service saving off the master service definition
- create the trigger file on the first slave pod
- wait 20 seconds for the failover to complete on the slave pod
- edit the slave pod's label to match that of the master
- recreate the master service using the stored service definition
- loop through the other remaining slave and delete its pod

At this point, clients when access the master's service will actually be accessing the new master. Also, Openshift will recreate the number of slaves to its original configuration which each slave pointed to the new master. Replication from the master to the new slaves will be started as each new slave is started by Openshift.

Openshift Example 13 - Metrics Collection

This example shows how postgres metrics can be collected and stored in prometheus and graphed with grafana.

First, create the crunchy-scope pod which contains the prometheus data store and the grafana graphing web application:

```
edit scope-pv.json to update the NFS share info
oc create -f scope-pv.json
oc create -f scope-pvc.json
oc process -f scope-nfs.json | oc create -f -
```

At this point, you can view the prometheus web console at crunchy-scope:9090, the prometheus push gateway at crunchy-scope:9091, and the grafana web app at crunchy-scope:3000.

Next, start a postgres pod that has the crunchy-collect container as follows:

```
oc process -f master-collect.json | oc create -f -
```

At this point, metrics will be collected every 3 minutes and pushed to prometheus. You can build graphs off the metrics using grafana.

Openshift Example 14 - Vacuum

This example shows how you can run a vacuum job against a postgres database container.

The crunchy-vacuum container image exists to allow a DBA a way to run a job either one-off or scheduled to perform a variety of vacuum operations.

To run the vacuum a single time, an example is included as follows from the examples/openshift directory:

```
cd examples/openshift/master-slave
./run.sh
cd ../vacuum-job
./run.sh
```

This will start a vacuum container that runs as a Kube Job type. It will run once. The crunchy-vacuum image is executed, passed in the Postgres connection parameters to the single-master postgres container. The type of vacuum performed is dictated by the environment variables passed into the job. The complete set of environment variables read by the vacuum job include:

- `VAC_FULL` - when set to true adds the `FULL` parameter to the `VACUUM` command
- `VAC_TABLE` - when set, allows you to specify a single table to vacuum, when not specified, the entire database tables are vacuumed
- `JOB_HOST` - required variable is the postgres host we connect to
- `PG_USER` - required variable is the postgres user we connect with
- `PG_DATABASE` - required variable is the postgres database we connect to
- `PG_PASSWORD` - required variable is the postgres user password we connect with
- `PG_PORT` - allows you to override the default value of 5432
- `VAC_ANALYZE` - when set to true adds the `ANALYZE` parameter to the `VACUUM` command
- `VAC_VERBOSE` - when set to true adds the `VERBOSE` parameter to the `VACUUM` command
- `VAC_FREEZE` - when set to true adds the `FREEZE` parameter to the `VACUUM` command

Openshift Example 15 - Custom Configuration Files

This example shows how you can use your own customized version of `setup.sql` when creating a postgres database container.

If you mount a `/pgconf` volume, crunchy-postgres will look at that directory for `postgresql.conf`, `pg_hba.conf`, and `setup.sql`. If it finds one of them it will use that file instead of the default files.

The example shows how a custom `setup.sql` file can be used. Run it as follows from the `examples/openshift/custom-config` directory:

```
./run.sh
```

This will start a database container that will use an NFS mounted /pgconf directory that will contain the custom setup.sql file found in the example directory.

Openshift Example 16 - pgbouncer

This example shows how you can use the crunchy-pgbouncer container when running under Openshift.

The example assumes you have run the master/slave example found here:

```
examples/openshift/master-slave-dc
```

Then you would start up the pgbouncer container using the following example:

```
examples/openshift/pgbouncer
```

The example assumes you have an NFS share path of /nfsfileshare/! NFS is required to mount the pgbouncer configuration files which are then mounted to /pgconf in the crunchy-pgbouncer container.

If you mount a /pgconf volume, crunchy-postgres will look at that directory for postgresql.conf, pg_hba.conf, and setup.sql. If it finds one of them it will use that file instead of the default files.

Test the example by killing off the master database container as follows:

```
oc delete pod pg-master-rc-dc
```

Then watch the pgbouncer log as follows to confirm it detects the loss of the master:

```
oc logs pgbouncer
```

After the failover is completed, you should be able to access the new master using the master service as follows:

```
psql -h pg-master-rc-dc.openshift.svc.cluster.local -U master postgres
```

and access the slave as follows:

```
psql -h pg-slave-rc-dc.openshift.svc.cluster.local -U master postgres
```

or via the pgbouncer proxy as follows:

```
psql -h pgbouncer.openshift.svc.cluster.local -U master master
```

Openshift Example 17 - synchronous slave

This example deploys a PostgreSQL cluster with a master, a synchronous slave, and an asynchronous slave. The two slaves share the same Service.

Running the example:

```
examples/openshift/sync/run.sh
```

Connect to the **master** and **slave** databases as follows:

```
psql -h master -U postgres postgres -c 'create table mister (id int)'  
psql -h master -U postgres postgres -c 'insert into mister values (1)'  
psql -h master -U postgres postgres -c 'table pg_stat_replication'  
psql -h slave -U postgres postgres -c 'select inet_server_addr(), * from mister'  
psql -h slave -U postgres postgres -c 'select inet_server_addr(), * from mister'  
psql -h slave -U postgres postgres -c 'select inet_server_addr(), * from mister'
```

This set of queries will show you the IP address of the Postgres slave container, notice it changes because of the round-robin Service proxy we are using for both slaves. The example queries also show that both slaves are replicating from the master.

Openshift Example 18 - pgadmin4

This example, `examples/openshift/pgadmin4`, provides a container that runs the pgadmin4 web application.

To run this example, run the following:

```
cd $BUILD_BASE/examples/openshift/pgadmin4  
./run.sh
```

This script creates the **pgadmin4** pod and service, it will expose port 5050.

You should now be able to browse to <http://pgadmin4.openshift.svc.cluster.local:5050> and log into the web application using a user ID of **admin@admin.org** and password of **password**. Replace YOURLOCALIP with whatever your local IP address happens to be.

Openshift Example 19 - workshop

This example, `examples/openshift/workshop`, provides an example of using Openshift Templates to build pods, routes, services, etc.

You use the **oc new-app** command to create objects from the JSON templates. This is an alternative way to create Openshift objects instead of using **oc create**.

This example is used within a joint Redhat-Crunchy workshop that is given at various conferences to demonstrate Openshift and Crunchy Containers working together. Thanks to Steven Pousty from Redhat for this example!

See the README file within the workshop directory for instructions on running the example.

Openshift Example 20 - PITR (point in time recovery)

This is a complex example. For details on how PITR is implemented within the Suite, see the [PITR Documentation](#) for details and background.

This example, `examples/openshift/pitr`, provides an example of performing a PITR using Openshift.

Lets start by running the example database container:

```
cd $BUILDBASE/examples/openshift/pitr
./run-master-pitr.sh
```

This step will create a database container, **master-pitr**. This container is configured to continuously write WAL segment files to a mounted volume (`/pgwal`). The mounted volume persists to an NFS persistent volume, **master-pitr-wal-pvc**. This container writes its data files to another NFS volume, **master-pitr-pvc**.

After you start the database, you will create a base backup using this command:

```
./run-master-pitr-backup.sh
```

This will create a backup and write the backup files to a persistent volume (`/pgbackup`) using the **backup-master-pitr-pvc** volume claim. This volume is also mounted on NFS.

Next, lets create some recovery targets within the database, run the SQL commands against the **master-pitr** database as follows:

```
./run-sql.sh
```

This will create recovery targets named **beforechanges**, **afterchanges**, and **nomorechanges**. It will create a table, **pitrtest**, between the **beforechanges** and **afterchanges** targets. It will also run a SQL CHECKPOINT to flush out the changes to WAL segments.

Next, now that we have a base backup and a set of WAL files containing our database changes, we can shut down the **master-pitr** database to simulate a database failure. Do this by running the following:

```
oc delete pod master-pitr
```

Next, we will create 3 different restored database containers based upon the base backup and the saved WAL files.

First, we restore prior to the **beforechanges** recovery target. This recovery point is **before** the **pitrtest** table is created.

Edit the master-pitr-restore.json file, and edit the environment variable to indicate we want to use the **beforechanges** recovery point:

```
, {  
  "name": "RECOVERY_TARGET_NAME",  
  "value": "beforechanges"  
}, {
```

Then run the following to create the restored database container:

```
./run-restore-pitr.sh
```

After the database has restored, you should be able to perform a test to see if the recovery worked as expected:

```
psql -h master-pitr-restore.openshift.svc.cluster.local -U postgres postgres -c 'table  
pitrtest'  
psql -h master-pitr-restore.openshift.svc.cluster.local -U postgres postgres -c  
'create table foo (id int)'  
psql -h master-pitr-restore.openshift.svc.cluster.local -U postgres postgres -c  
'select pg_xlog_replay_resume()'  
psql -h master-pitr-restore.openshift.svc.cluster.local -U postgres postgres -c  
'create table foo (id int)'
```

The output of these command should show that the **pitrtest** table is not present. It should also show that you can not create a new table because the database is paused in recovery mode. Lastly, if you execute a **resume** command, it will show that you can now create a table as the database has fully recovered.

You can also test that if **afterchanges** is specified, that the **pitrtest** table is present but that the database is still in recovery mode.

Lastly, you can test a full recovery using **all** of the WAL files, if you remove the **RECOVERY_TARGET_NAME** environment variable completely.

The NFS portions of this example depend upon an NFS file system with the following path configurations be present:

```
/nfsfileshare  
/nfsfileshare/backups  
/nfsfileshare/WAL
```

Openshift Tips

Tip 1: Finding the Postgresql Passwords

The passwords used for the PostgreSQL user accounts are generated by the Openshift 'process' command. To inspect what value was supplied, you can inspect the master pod as follows:

```
oc get pod pg-master-rc-1-n5z8r -o json
```

Look for the values of the environment variables: - PG_USER - PG_PASSWORD - PG_DATABASE

Tip 2: Examining a backup job log

Database backups are implemented as a Kubernetes Job. A Job is meant to run one time only and not be restarted by Kubernetes. To view jobs in Openshift you enter:

```
oc get jobs  
oc describe job backupjob
```

You can get detailed logs by referring to the pod identifier in the job 'describe' output as follows:

```
oc logs backupjob-pxh2o
```

Tip 3: Backup Lifecycle

Backups require the use of network storage like NFS in Openshift. There is a required order of using NFS volumes in the manner we do database backups.

So, first off, there is a one-to-one relationship between a PV (persistent volume) and a PVC (persistence volume claim). You can NOT have a one-to-many relationship between PV and PVC(s).

So, to do a database backup repeatably, you will need to following this general pattern: * as openshift admin user, create a unique PV (e.g. backup-pv-mydatabase) * as a project user, create a unique PVC (e.g. backup-pvc-mydatabase) * reference the unique PVC within the backup-job template * execute the backup job template * as a project user, delete the job * as a project user, delete the pvc * as openshift admin user, delete the unique PV

This procedure will need to be scripted and executed by the devops team when performing a database backup.

Tip 4: Persistent Volume Matching

Restoring a database from an NFS backup requires the building of a PV which maps to the NFS backup archive path. For example, if you have a backup at `/backups/pg-foo/2016-01-29:22:34:20` then we create a PV that maps to that NFS path. We also use a "label" on the PV so that the specific backup PV can be identified.

We use the pod name in the label value to make the PV unique. This way, the related PVC can find the right PV to map to and not some other PV. In the PVC, we specify the same "label" which lets Kubernetes match to the correct PV.

Tip 5: Restore Lifecycle

To perform a database restore, we do the following: * locate the NFS path to the database backup we want to restore with * edit a PV to use that NFS path * edit a PV to specify a unique label * create the PV * edit a PVC to use the previously created PV, specifying the same label used in the PV * edit a database template, specifying the PVC to be used for mounting to the `/backup` directory in the database pod * create the database pod

If the `/pgdata` directory is blank AND the `/backup` directory contains a valid postgres backup, it is assumed the user wants to perform a database restore.

The restore logic will copy `/backup` files to `/pgdata` before starting the database. It will take time for the copying of the files to occur since this might be a large amount of data and the volumes might be on slow networks. You can view the logs of the database pod to measure the copy progress.

Tip 6: Password Mgmt

Remember that if you do a database restore, you will get whatever user IDs and passwords that were saved in the backup. So, if you do a restore to a new database and use generated passwords, the new passwords will not be the same as the passwords stored in the backup!

You have various options to deal with managing your passwords.

- externalize your passwords using secrets instead of using generated values
- manually update your passwords to your known values after a restore

Note that you can edit the environment variables when there is a 'dc' using, currently only the slaves have a 'dc' to avoid the possibility of creating multiple masters, this might need to change in the future, to better support password management:

```
oc env dc/pg-master-rc PG_MASTER_PASSWORD=foo PG_MASTER=user1
```

Tip 7: Log Aggregation

Openshift can be configured to include the EFK stack for log aggregation. Openshift Administrators can configure the EFK stack as documented here:

https://docs.openshift.com/enterprise/3.1/install_config/aggregate_logging.html

Tip 8: nss_wrapper

If an Openshift deployment requires that random generated UUIDs be supported by containers, the Crunchy containers can be modified similar to those located here to support the use of nss_wrapper to equate the random generated UUIDs/GIDs by openshift with the postgres user:

<https://github.com/openshift/postgresql/blob/master/9.4/root/usr/share/container-scripts/postgresql/common.sh>

Tip 9: build box setup

golang is required to build the pgbadger container, on RH 7.2, golang is found in the 'server optional' repository and needs to be enabled to install.

golang is required to build the pgbadger container, on RH 7.2, golang is found in the 'server optional' repository and needs to be enabled to install.

Tip 10: encoding secrets

You can use kubernetes secrets to set and maintain your database credentials. Secrets requires you base64 encode your user and password values as follows:

```
echo -n 'myuserid' | base64
```

You will paste these values into your JSON secrets files for values.

docker to be installed.

You can keep yum from upgrading docker by including this line in your /etc/yum.conf file:

```
exclude=docker-1.9* docker-selinux-1.9*
```

Tip 12: DNS configuration for Openshift development

Luke Meyer from Redhat wrote an excellent blog on how to configure dnsmasq and Openshift, it is located here:

<http://developers.redhat.com/blog/2015/11/19/dns-your-openshift-v3-cluster/>

Key things included in this blog are:

- configuring dhcp to include the local IP address in /etc/resolv.conf upon boot
- configuring dnsmasq
- configuring openshift dns to listen on another port

In my dev setup, I have openshifts DNS listening on 127.0.0.1:8053. I have my dnsmasq listening on the local IP address 192.168.0.109:53

Therefore in my /etc/dhcp/dhclient.conf I have this config:

```
prepend domain-name-servers 192.168.0.109;
```

If you dont have your DNS configured correctly, replication controllers and deployment configs basically will not work.

Legal Notices

Copyright © 2016 Crunchy Data Solutions, Inc.

CRUNCHY DATA SOLUTIONS, INC. PROVIDES THIS GUIDE "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF NON INFRINGEMENT, MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Crunchy, Crunchy Data Solutions, Inc. and the Crunchy Hippo Logo are trademarks of Crunchy Data Solutions, Inc.